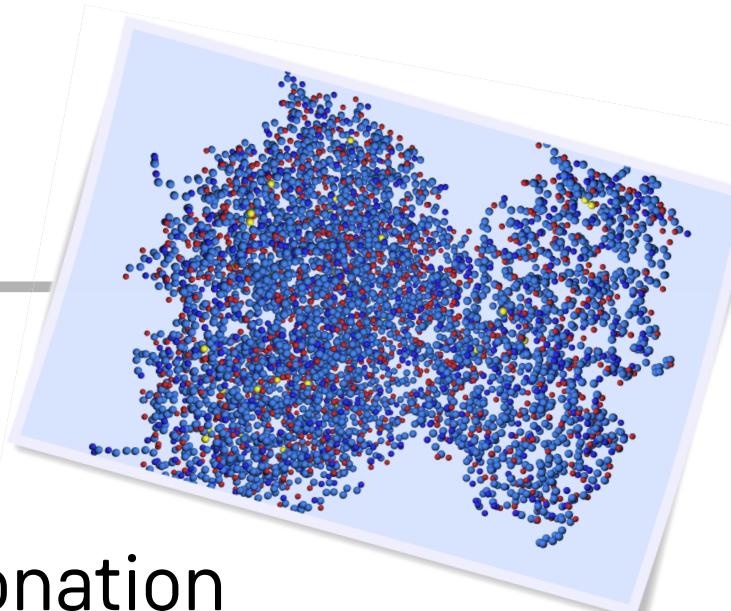


Make HPC Software Great Again

Developing A Fast Multipole Toolbox For Modern Applications On Modern Hardware

A. Beckmann, L. Morgenstern, I. Kabadshow | Jülich Supercomputing Centre, Research Centre Jülich, Germany
 C. Kutzner, Th. Ullmann, B. Kohnke | Max Planck Institute for Biophysical Chemistry, Göttingen, Germany

① Scientific Background



- ❖ **DFG-funded Project:** GROMEX (2013-2019)
Unified long-range electrostatics and dynamic protonation for realistic biomolecular simulations on the Exascale

- ❖ **Application:** Molecular Dynamics Code GROMACS

- ❖ **Library:** Fast Multipole Method (FMM) with $O(N)$ complexity



① Starting Point

- ✓ GROMACS is written in C++11
- ❶ GROMACS is tightly coupled with PME as long range solver
- ❷ No external libraries are used for long range interactions until now



- ❸ FMM available in highly optimized Fortran90 version only
- ❹ 150k+ lines of code, optimized for minimal FLOPs
- ❺ Supports only a few platforms via hand-written C-intrinsics
- ❻ Code duplications, non-generic codebase, no threading support

② Challenges

Hardware flexibility

- ILP, SIMD, OoOE
- Cache levels & sizes
- NUMA
- Threading
- Message passing

Lifecycle of hardware

- Changes every year

Algorithmic configurability

- Different implementations of compute kernels
- Different critical paths

Lifecycle of staff

- PhD every 3-5 years
- Intern every 6-12 months

Application customization

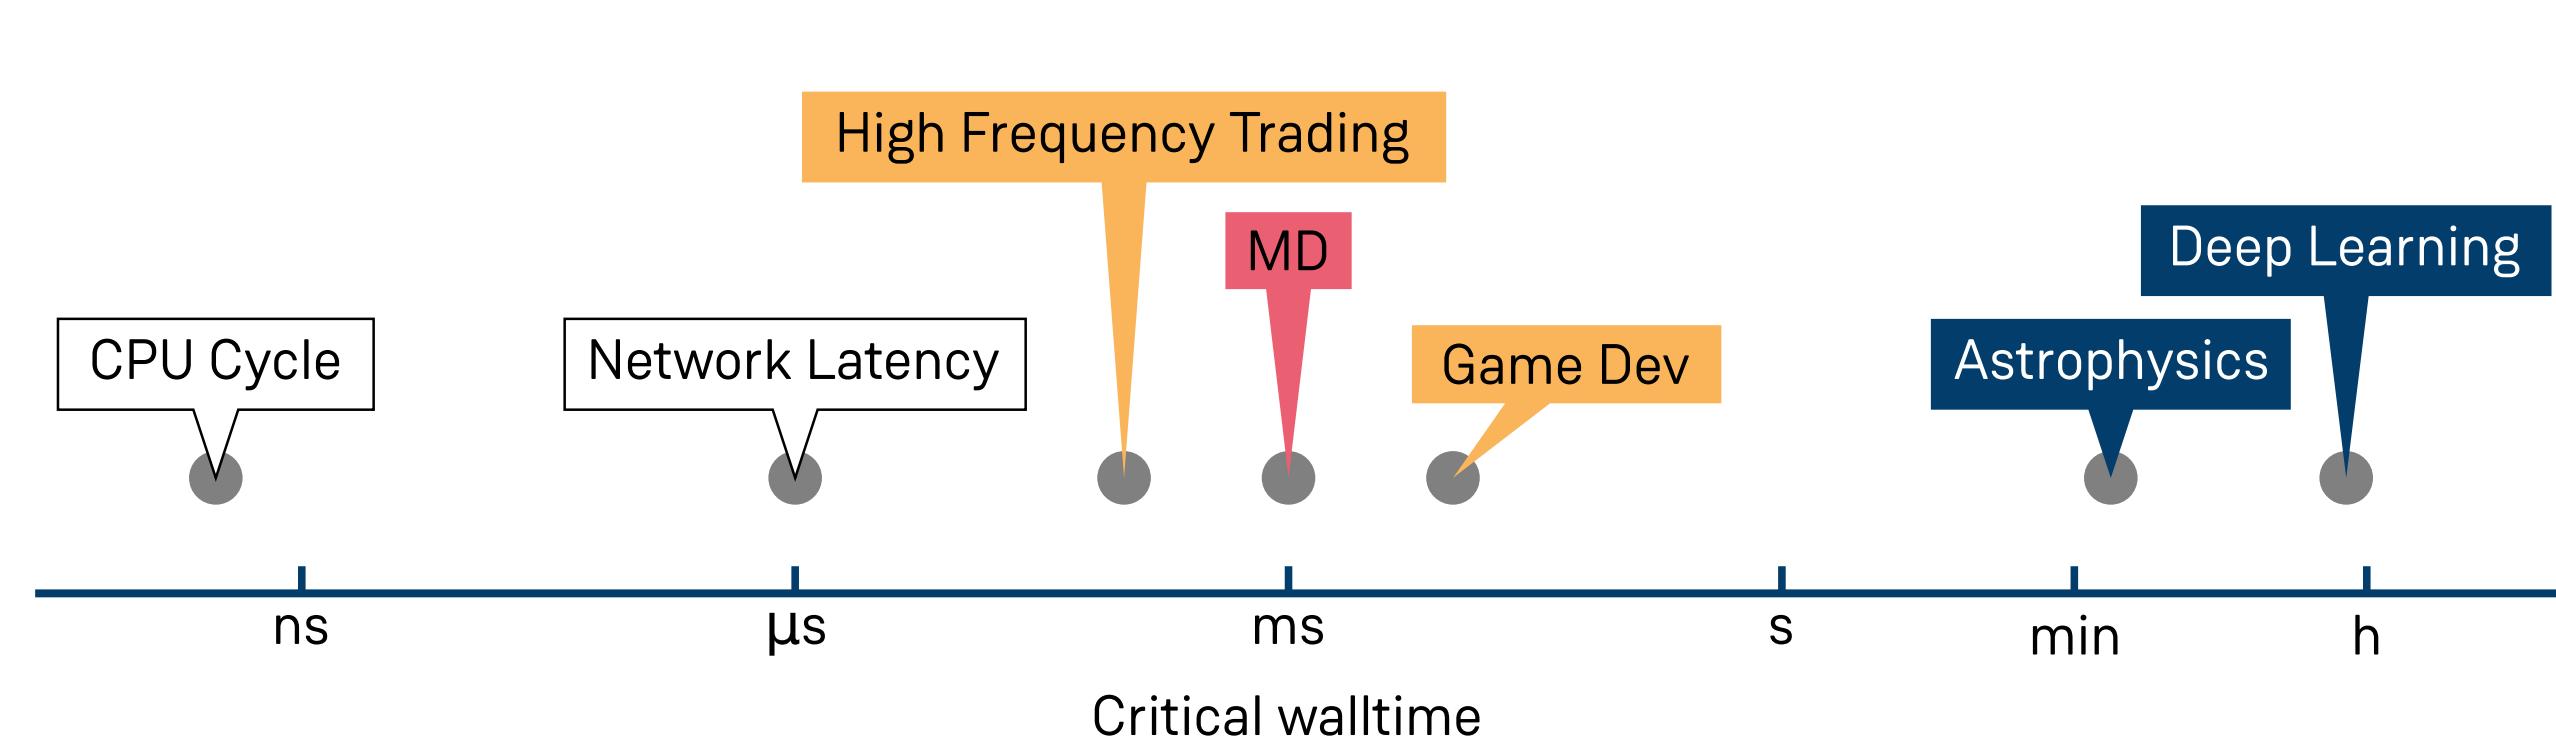
- Physical model
- Accuracy range
- System size

Lifecycle of code

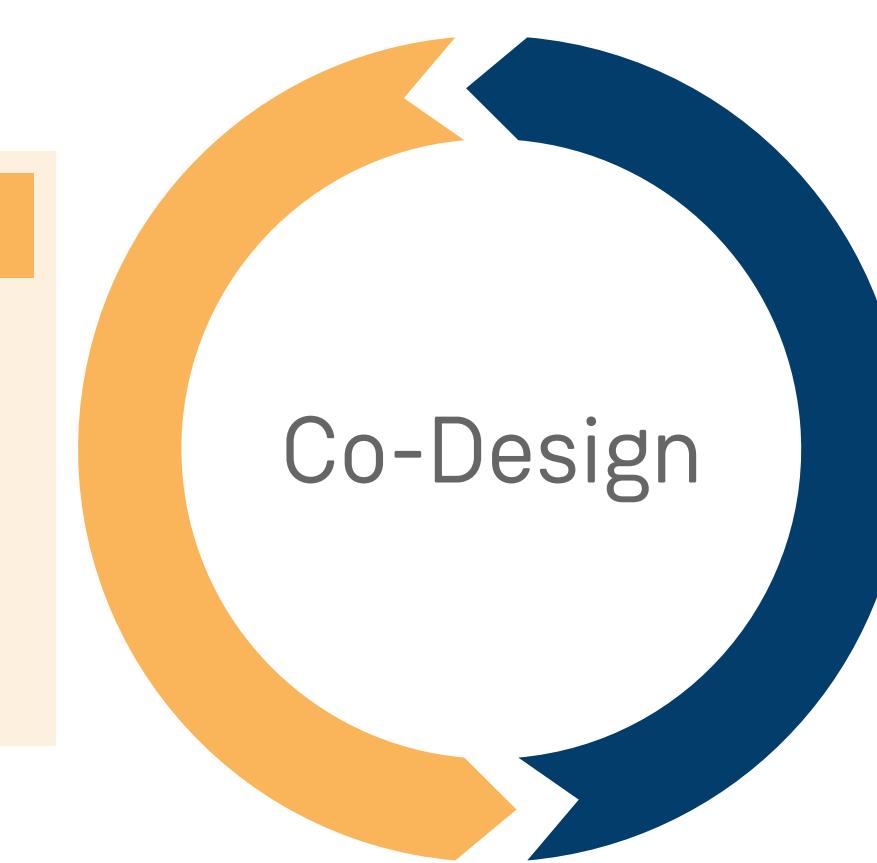
- Minutes to lifetime

③ HPC≠HPC

- Non-trivial data structures
- Latency-critical walltime
- Non-uniform parallelism
- ❶ Parallelism on all levels required



④ Co-Design Approach



Application developer

- Algorithmic details
- Interface design

Library developer A

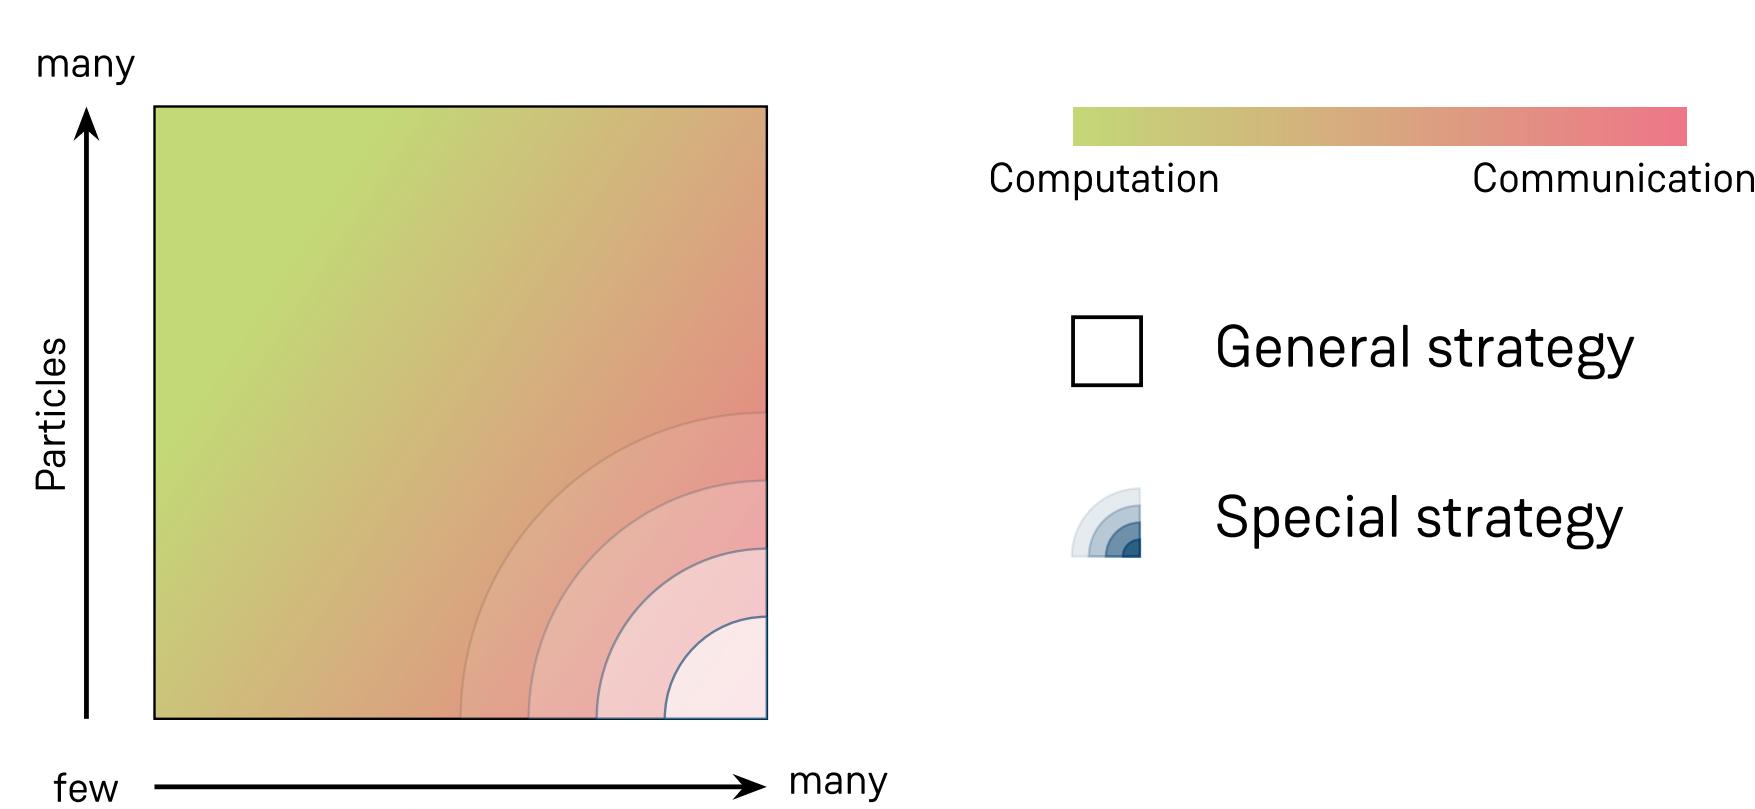
- HW knowledge
- Interface design
- Software engineering

Library developer B

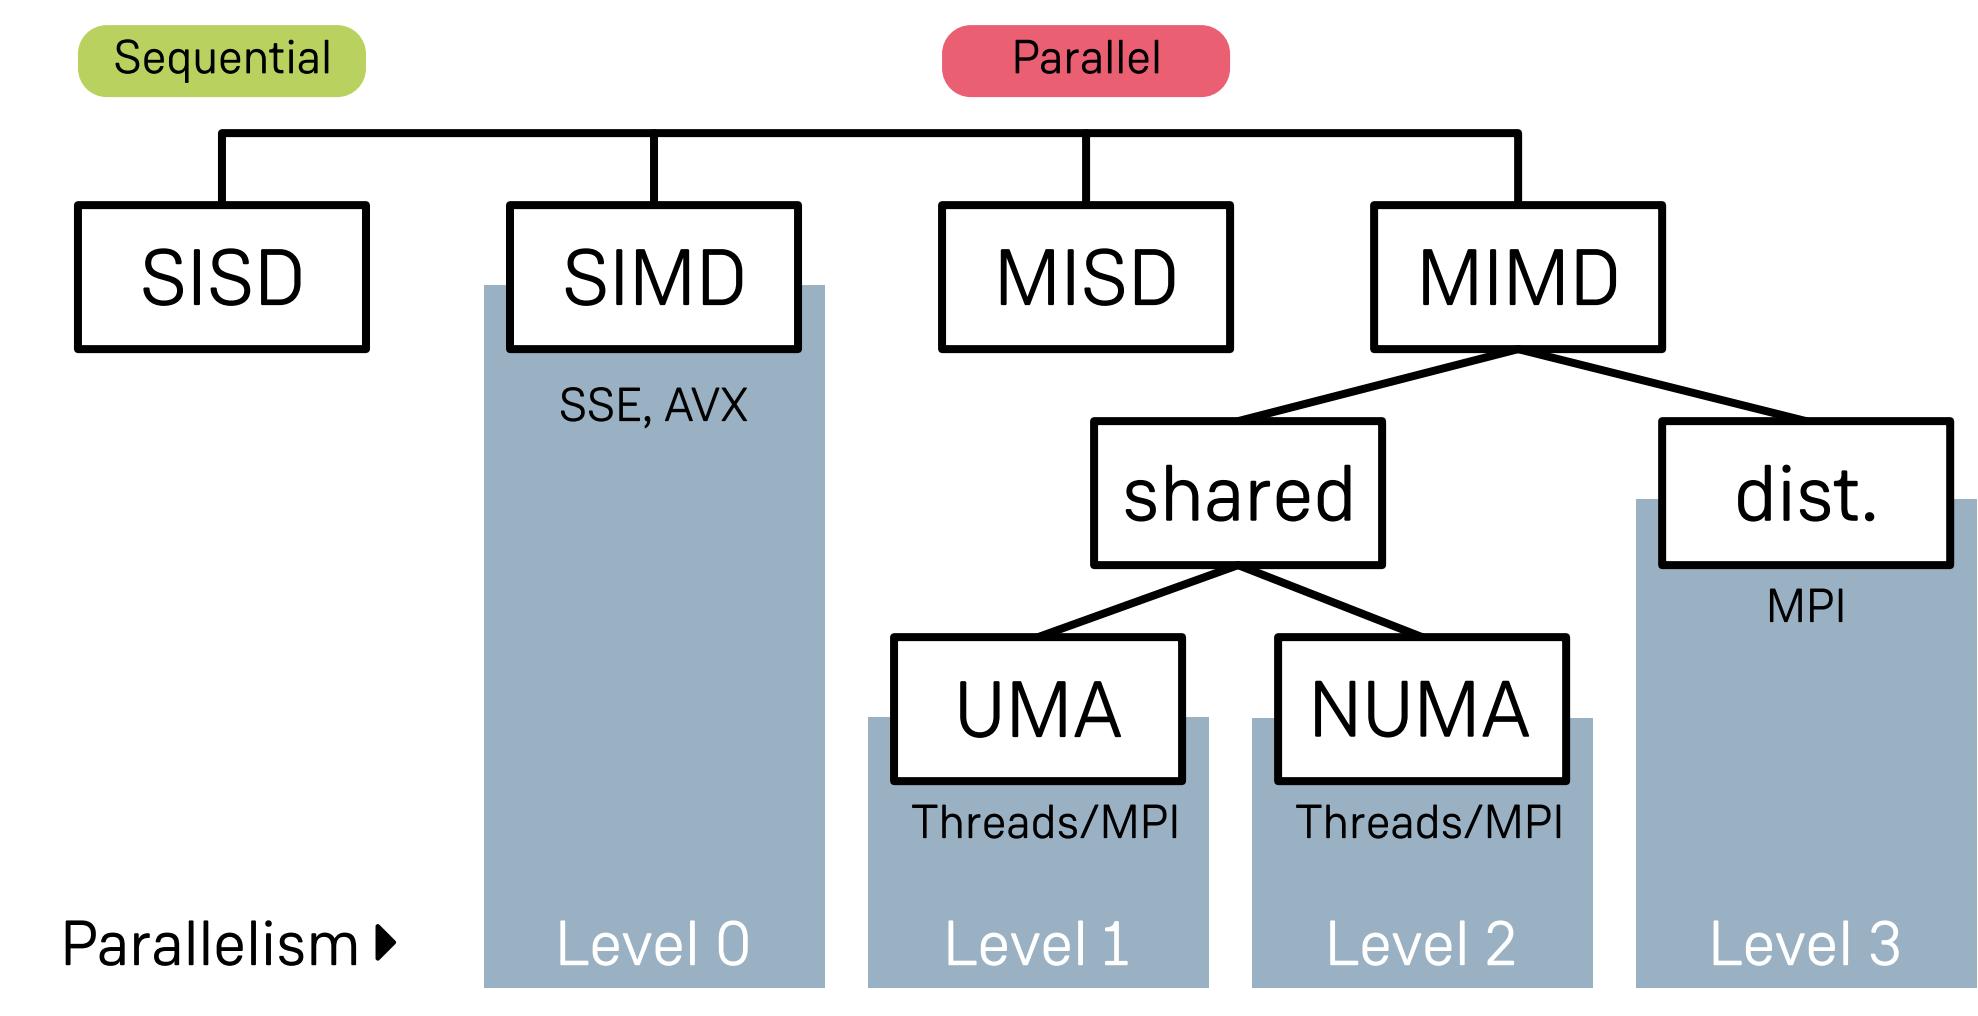
- Extensions
- Unit tests

⑥ Algorithm-Aware Specialization

- Parallelization strategy depends on the critical path of the algorithm
- Strategy must be interchangeable without 'large' changes to the user code
- Algorithm knowledge allows for better performance



⑤ Modularity of Parallelism



- Decouple parallelization strategies from algorithmic workflow
- Different parallelization levels designed as independent sub-libraries, e.g. vectorization, tasking

⑦ Conclusion

- Mature programming language like C++11 mandatory
- Template metaprogramming helps to write readable and performant code
- Encapsulation of features helps maintainability
- Lines of code reduced to 25k+ due to deduplication
- No performance drawback with respect to C++ classes
- Zero-overhead abstraction layers via templates

- Domain scientist deals with domain science not the hardware
- Computer scientist deals with software design & hardware only
- Co-design cycle of interfaces defines overlap of domain and computer scientist

- Reuse of individual components possible through library design
- Independent projects can benefit by using such libraries

